

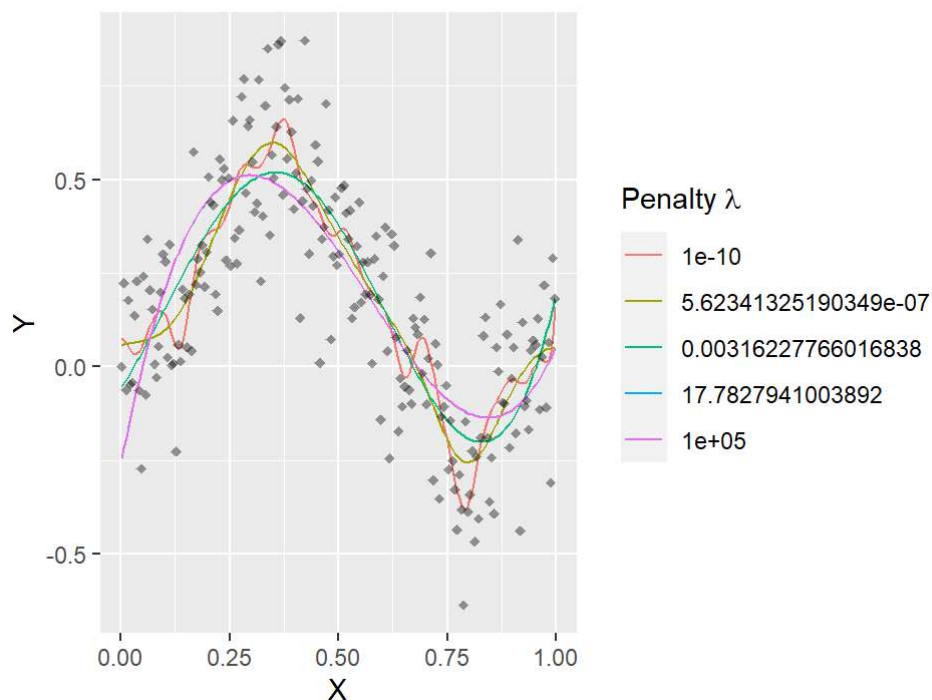
# HW5

Sili Fan and Kwan Ho Lee

6/1/2021

## Problem 1

Example cubic penalized regression spline with 30 knots




(a) Implement the cross-validation and generalized cross-validation methods for choosing the smoothing parameter  $\lambda$

Cross validation (CV)

We are interested in finding the value of  $\lambda$  that minimizes the following function for a particular sample  $\{x_i, y_i\}$ . We accomplish this with a simple grid search over an appropriate range of  $\lambda > 0$ .

$$CV(\lambda) = \frac{1}{N} \sum_{i=1}^N \left( \frac{y_i - \hat{f}_\lambda(x_i)}{1 - (S_\lambda)_{i,i}} \right)^2$$




Note that the minimizing smoothing penalty is  $\lambda_{CV} = 4.0370173 \times 10^{-8}$  with a cross validation score of  $CV(\lambda_{CV}) = 0.0310471$ .

### Generalized cross validation (GCV)

We are interested in finding the value of  $\lambda$  that minimizes the following function for a particular sample  $\{x_i, y_i\}$ . We accomplish this with a simple grid search over an appropriate range of  $\lambda > 0$ .


$$GCV(\lambda) = \frac{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{f}_\lambda(x_i))^2}{\left(1 - \frac{1}{N} \text{tr}(S_\lambda)\right)^2}$$



Note that the minimizing smoothing penalty is  $\lambda_{GCV} = 4.6415888 \times 10^{-8}$  with a generalized cross validation score of  $GCV(\lambda_{GCV}) = 0.0310048$ .

(b)

$$AIC_C(\lambda) = \log(\|y - \hat{f}_\lambda\|_2^2) + \frac{2(\text{tr}(H_\lambda) + 1)}{n - \text{tr}(H_\lambda) - 2}$$



Note that the minimizing smoothing penalty is  $\lambda_{AIC_C} = 5.7223677 \times 10^{-8}$  with a score of  $AIC_C(\lambda_{AIC_C}) = -3.4554865$

(c)

Proof:

$E(\|y - \hat{f}_\lambda\|^2) = \|(I - H_\lambda)f\|^2 + \sigma^2(tr(H_\lambda H_\lambda^\top) - 2tr(H_\lambda) + n)$  to get a sense for what this estimate will look like.

1.  $E(\|y - \hat{f}_\lambda\|^2) = E(\|(I - H_\lambda)y\|^2)$  because  $\hat{f}_\lambda = H_\lambda y$  and real matrices have both left and right distributivity.
2.  $E(\|(I - H_\lambda)y\|^2) = E(\|(I - H_\lambda)(f + e)\|^2)$  because  $y = f + e$ .
3.  $E(\|(I - H_\lambda)(f + e)\|^2) = E(\|(I - H_\lambda)f\|^2 + \|(I - H_\lambda)e\|^2)$  because  $E(e) = 0$  by assumption and the distributivity property noted in (1).
4.  $E(\|(I - H_\lambda)f\|^2 + \|(I - H_\lambda)e\|^2) = \|(I - H_\lambda)f\|^2 + E(\|(I - H_\lambda)e\|^2)$  because the lefthand quantity is a constant. It remains to show that  $E(\|(I - H_\lambda)e\|^2) = \sigma^2(tr(H_\lambda H_\lambda^\top) - 2tr(H_\lambda) + n)$
5.  $E(\|(I - H_\lambda)e\|^2) = \sigma^2 tr(I(H_\lambda - I)(H_\lambda - I)^\top I^\top)$  because  $Cov(e) = I\sigma^2$ .
6.  $tr(I(H_\lambda - I)(H_\lambda - I)^\top I^\top) = \sum_{i=1}^n (h_{ii} - 1)^2$ .
7.  $\sum_{i=1}^n (h_{\lambda,ii} - 1)^2 = tr(H_\lambda H_\lambda^\top) - 2tr(H_\lambda) + n$

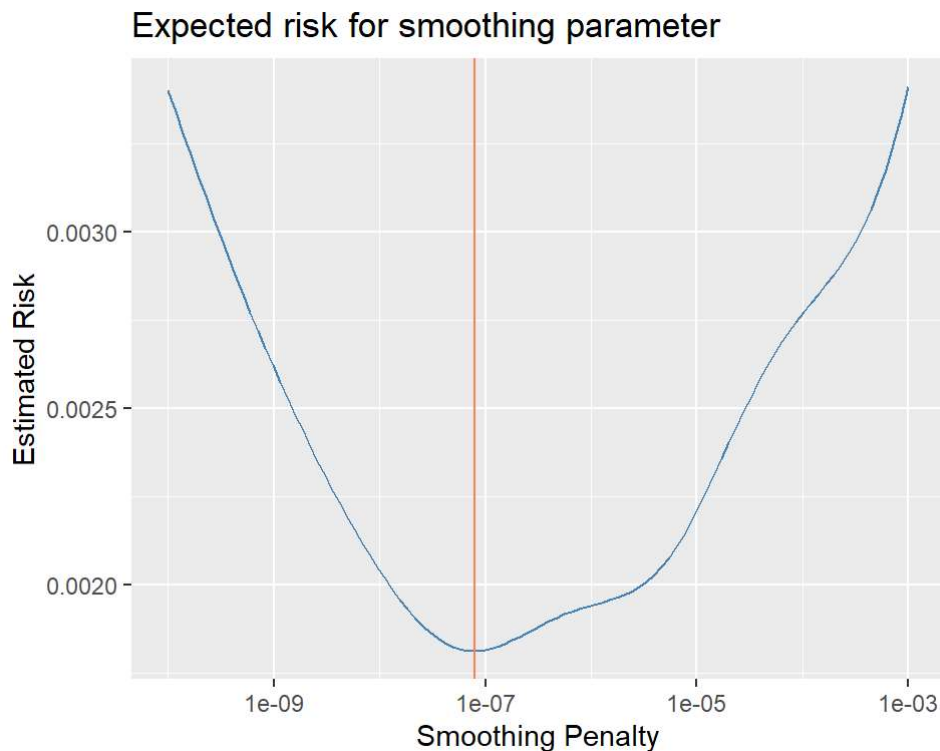
To complete the proof, we get the following steps, we develop an estimator for  $risk(\lambda) = E(\|f - \hat{f}_\lambda\|^2)$ :

1.  $risk(\lambda) = E(\|y - e - \hat{f}_\lambda\|^2)$  by the model assumption  $y = f + e$ .
2. Expanding the  $L_2$  norm in the expression above, using the linearity of expectation, and noting the fact that  $E(e) = 0$ ;  $E(\|y - e - \hat{f}_\lambda\|^2) = E(y^\top y - y^\top \hat{f}_\lambda + e^\top e - \hat{f}_\lambda^\top y + \hat{f}_\lambda^\top \hat{f}_\lambda)$
3. Consolidating terms in (2),  $risk(\lambda) = n\sigma^2 + E(\|y - \hat{f}_\lambda\|^2)$  and we use the method of moments to create an unbiased estimator when  $\sigma^2$  is known.

$y$  are the noisy observations,  $\hat{f}_\lambda = H_\lambda y$  where  $H_\lambda$  depends only on  $X$ , and we can create an estimator for  $\sigma^2$ :  $\hat{\sigma}^2 \approx RSS(\hat{f}_\lambda) = \frac{1}{n} \|y - \hat{f}_\lambda\|^2$ . Using this estimator for  $\sigma^2$ , the estimator for  $E(\|f - \hat{f}_\lambda\|^2)$  is proportional to  $\|y - \hat{f}_\lambda\|^2$ .

An alternative estimator for the risk, using classical pilots, is given in section 2.2.1 of Lee (2003): use  $\hat{f}_{\lambda_p} \approx f$  where  $\lambda_p$  is the  $CV$  parameter. This is the method that we employ in our algorithm.

The following algorithm exhibits this risk estimator minimization criterion for calculating the penalty parameter  $\lambda$  for a single experimental simulation parameterization.



Note that the minimizing smoothing penalty is  $\lambda_{ER} = 7.924829 \times 10^{-8}$  with a score of  $ER(\lambda_{ER}) = 0.0018119$ .

(d) Conduct a simulation study to compare the above four methods for choosing  $\lambda$ . Use the experimental setting from the paper downloadable from Canvas.

In this homework, we simulate six different methods for selecting a penalty parameter  $\lambda$  under six different noise, design density, spatial variation, and variance heterogeneity states. The four methods for fitting the penalty parameter  $\lambda$  for a penalized cubic regression spline are minimization of the following scores: cross validation, generalized cross validation, corrected Akaike information criterion, and expected risk.

In this part, we illustrate the concepts required for the simulation in the context of a simple example:  $y = f(x) + \epsilon$ .

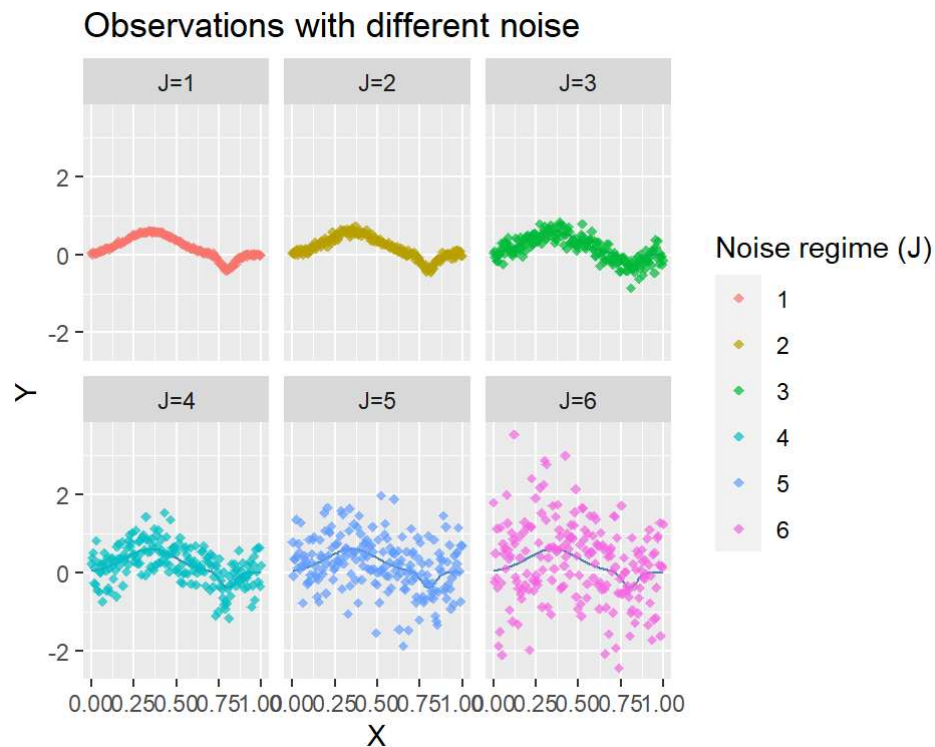
To fit a cubic penalized regression spline  $\hat{f}_\lambda(x)$  which approximates  $f(x)$ , we minimize the following functional within the class  $\mathbf{F}$  of cubic regression splines:

$$\hat{f}_\lambda(x) = \operatorname{argmin}_{f \in \mathbf{F}} \sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \sum_{j=1}^k \beta_{p_j}^2 = \|y - f(x)\|_2^2 + \lambda \beta^\top D \beta$$

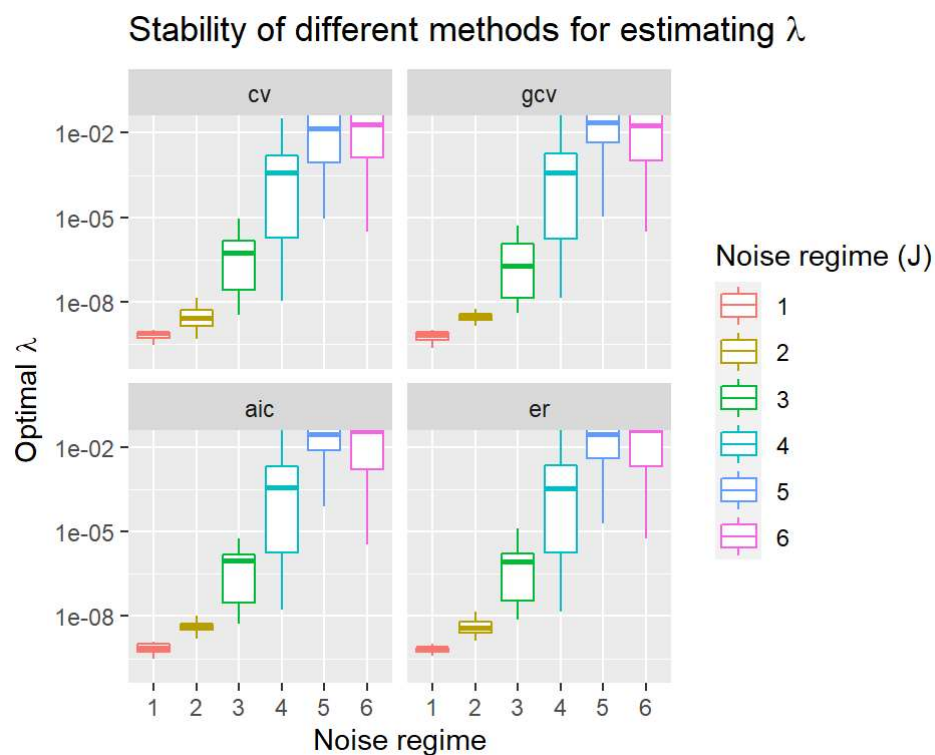
where  $D$  is the diagonal matrix with  $p + 1$  0s followed by  $K$  1s.

### Performance under different levels of noise

We evaluate the performance of  $\lambda$  estimators under the model  $y_{ij} = f(x_i) + \sigma_j \epsilon_i$  where  $\sigma_j = 0.02 + 0.04(j - 1)^2$  and  $\epsilon_i \sim \text{iid}N(0, 1)$ .



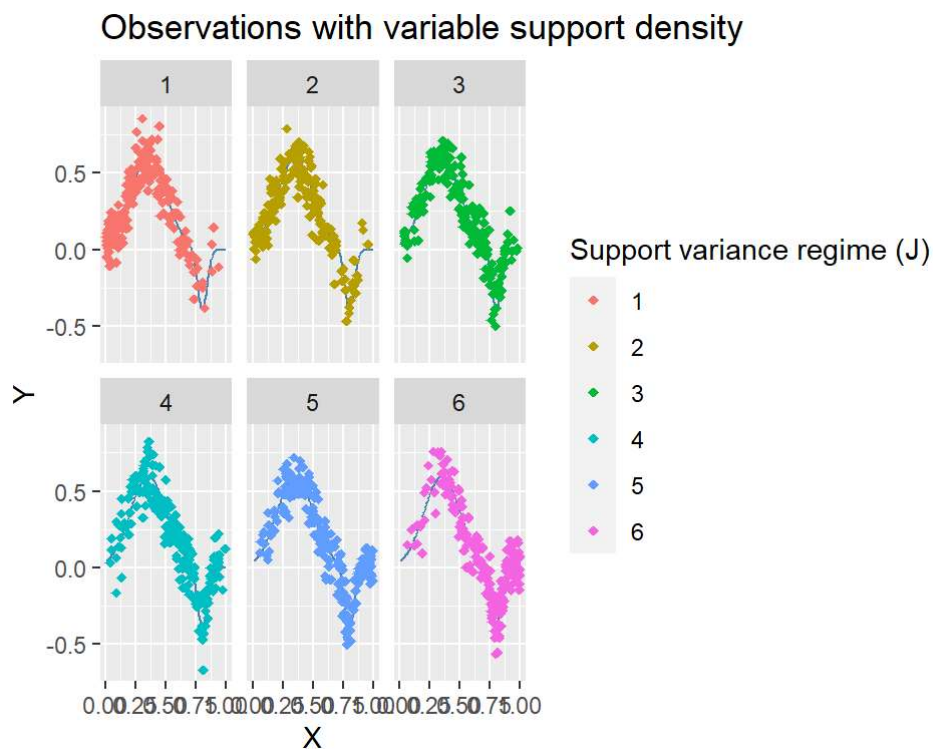
The above plot illustrates the effects of noise regimes  $j \in \{1, \dots, J = 6\}$ .



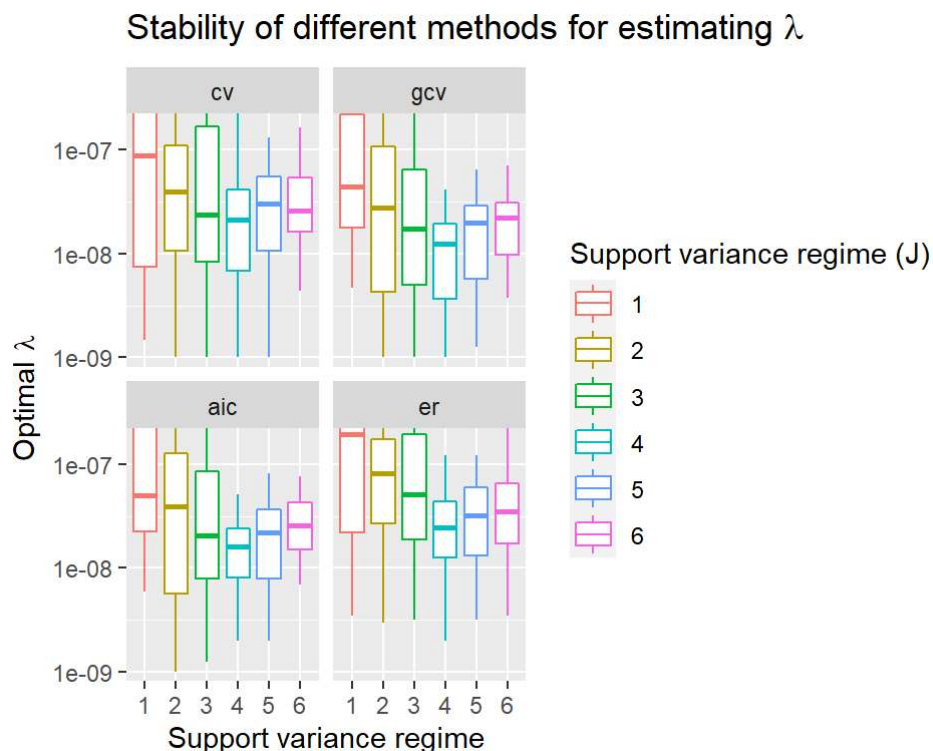
As the level of noise increases, the optimal penalty parameter increases and its stability decreases. The increase in  $\lambda$  means that we encourage the spline to be smoother - to fit less tightly to the noisy data - as the level of noise increases.

### Performance under different design density

$$y_{ij} = f(X_{ij}) + \sigma \epsilon_i \quad \text{where} \quad \sigma = 0.1, X_{ij} = F^{-1}(X_i), F(X) = \text{Beta}\left(\frac{j+4}{5}, \frac{11-j}{5}\right), X_i \sim \text{Unif}(0, 1)$$



As illustrated above, the small values of  $j$  group points to the left-hand side of the support, while the large values of  $j$  group points to the right-hand side.

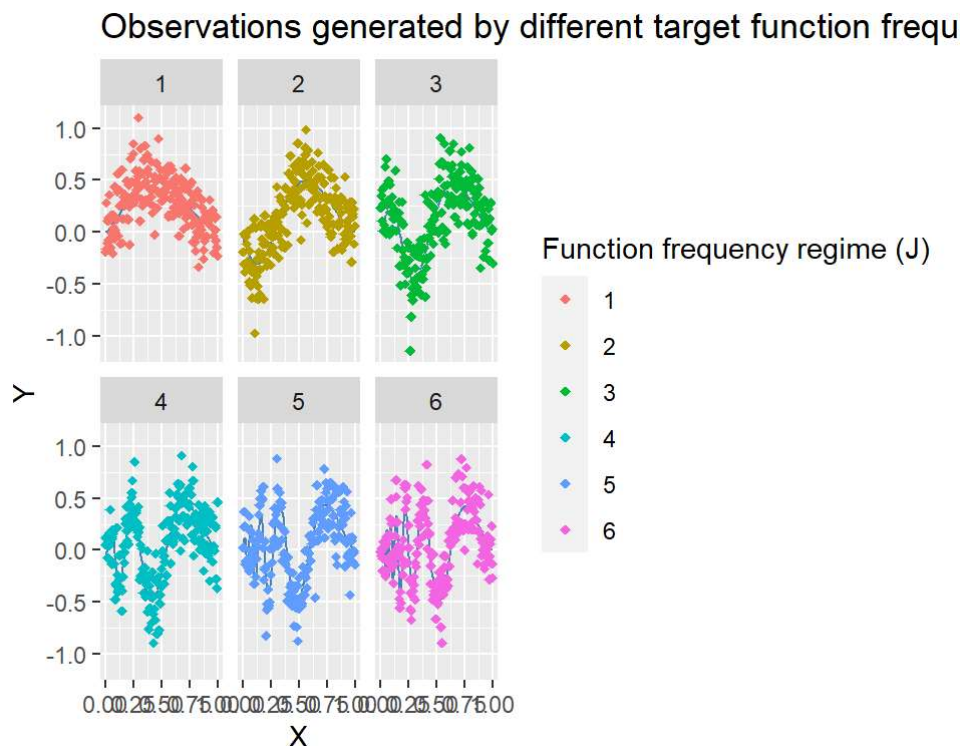


The stability of the penalty parameter estimations is much lower when the points are not homogeneously dense throughout the support. This is especially true considering that we fitting a spline with 30 equally spaced knots. A variable knot selection protocol might mitigate some of the instability of the penalty parameter estimates observed for  $j \in \{1, 2, 5, 6\}$ .

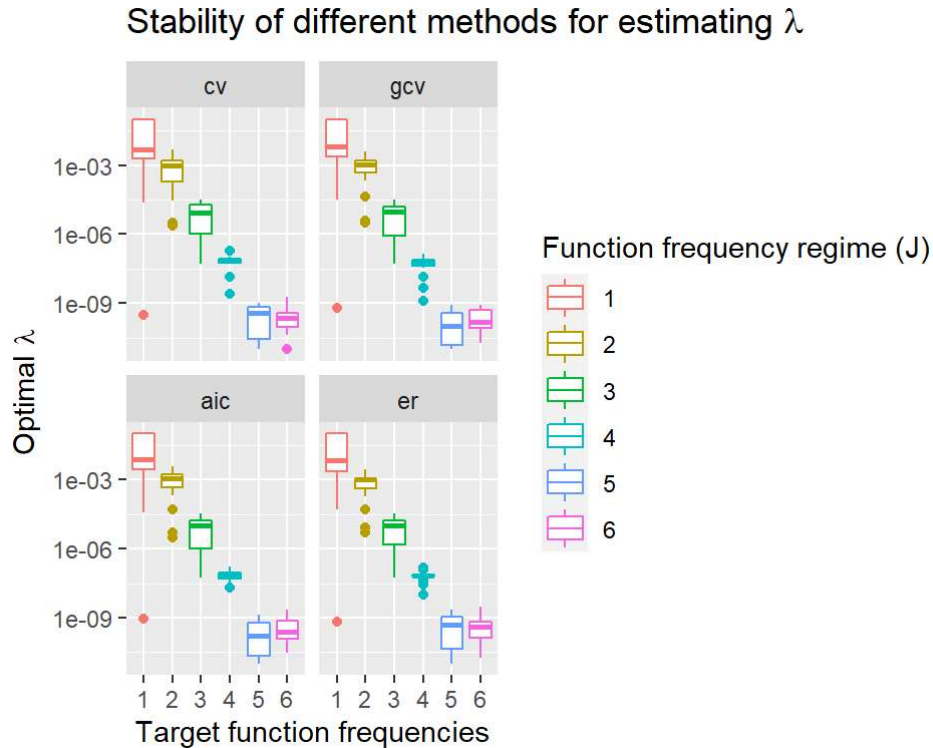
### Performance under different frequency of target function



$$y_{ij} = f_j(x_i) + \sigma\epsilon_i \quad \text{where} \quad \sigma = 0.2, f_j(x) = \sqrt{x * (1 - x)} \sin\left(\frac{2\pi(1 + 2^{(9-4j)/5})}{x + 2^{(9-4j)/5}}\right), \epsilon_i \sim \text{iid}N(0, 1)$$



As illustrated above, increasing  $j$  increases the frequency of the spatial distribution function  $f(x)$ .



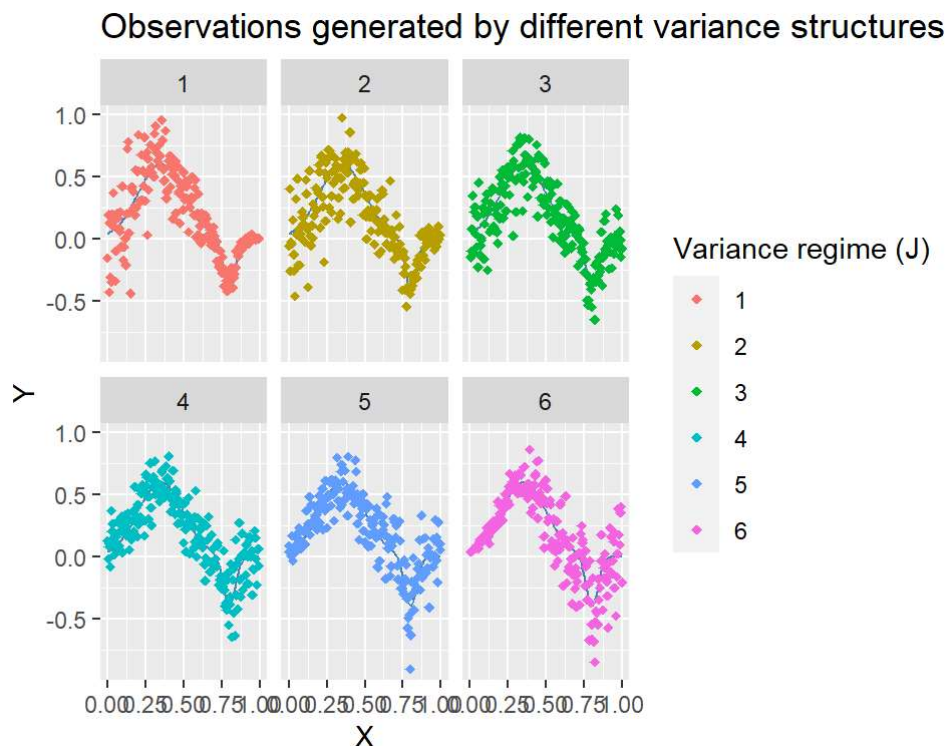
Increased spatial frequency stabilizes the estimates of the penalty parameter to a point. Furthermore, as the function becomes more erratic, the value of the penalty parameter  $\lambda$  decreases, forcing the spline to hug the data more closely.

### Performance under heterogeneous variances

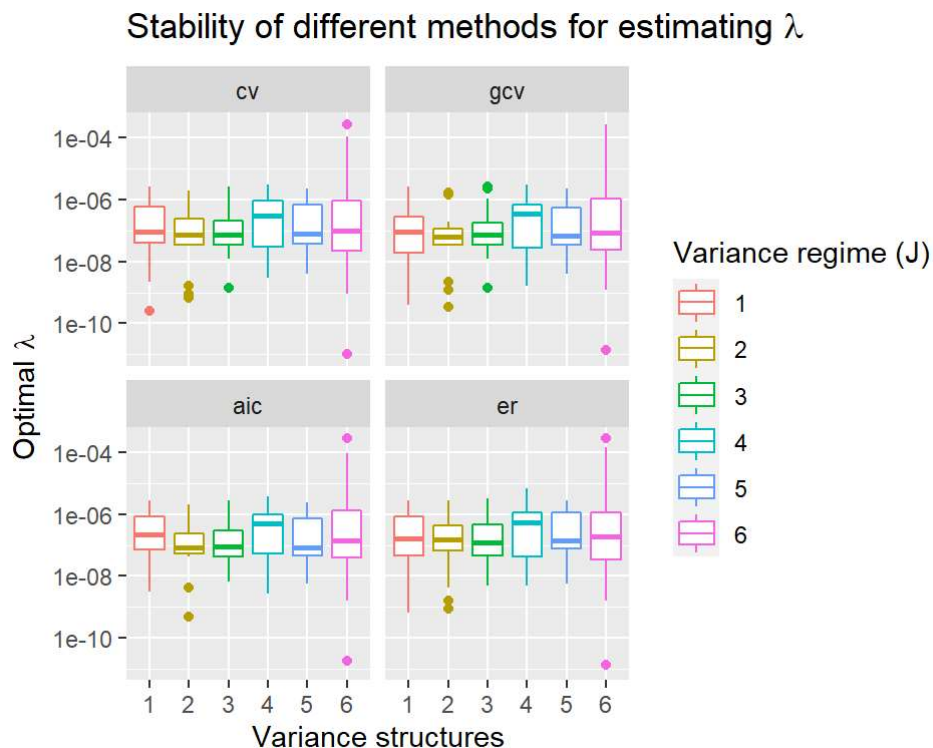


$$y_{ij} = f(x_i) + \sqrt{v_j(x_i)}\epsilon_i \quad \text{where} \quad v_j(x) = (0.15(1 + 0.4(2j - 7)(x - 0.5)))^2$$

and  $f(x)$  and  $\epsilon_i$  are defined as in the above.



As illustrated above, this method for varying the variance spatially incurs a high variance for the left-hand side when  $j$  is small and a high variance for the right-hand side when  $j$  is large.



Spatial variance modulation destabilizes smoothing parameter estimates. Having 30 equi-spaced knots is a serious issue here and a more dynamic collocation procedure would ameliorate some of this instability in the  $\lambda$  estimates.

# Appendix

The following code is used in Problem 1 part (a)

```
cv = function(l, k=k, X=X, Y=Y){  
  
  n = dim(X)[1]  
  D = diag(c(rep(0, dim(X)[2]-k), rep(1, k)))  
  H = X %*% solve(t(X) %*% X + l*D) %*% t(X)  
  fhys = H %*% Y  
  hii = diag(H)  
  r = mean(((Y-fhys)/(1-hii))^2)  
  return(r)  
}  
  
cvv = Vectorize(cv, vectorize.args = c("l"))  
lambdas = 10^(seq(-10, -4, length.out=100))  
cvs = cvv(X=X, Y=Y, l=lambdas, k=k)  
  
cv_df = data.frame(x=lambdas, y=cvs)  
plt_cv = ggplot(data=cv_df) +  
  geom_line(aes(x=x, y=y), color="steelblue") +  
  labs(x="Smoothing Penalty", y="Cross Validation Score",  
       title="Cross validated values of lambda") +  
  scale_x_log10() +  
  # geom_hline(yintercept = min(cvs), color="coral") +  
  geom_vline(xintercept = lambdas[which(cvs == min(cvs))], color="coral")
```

```
gcv = function(X=X, Y=Y, k=k, l){  
  
  n = dim(X)[1]  
  D = diag(c(rep(0, dim(X)[2]-k), rep(1, k)))  
  H = X %%% solve(t(X) %%% X + l*D) %%% t(X)  
  
  fhys = H %%% Y  
  tr = diag(H)  
  
  numer = mean((Y-fhys)^2)  
  denom = (1-mean(tr))^2  
  r = numer/denom  
  return(r)  
}  
  
gcvv = Vectorize(gcv, vectorize.args = c("l"))  
lambdas = 10^(seq(-9, -4, length.out=100))  
gcvs = gcvv(X=X, Y=Y, l=lambdas, k=k)  
  
gcv_df = data.frame(x=lambdas, y=gcvs)  
plt_gcv = ggplot(data=gcv_df) +  
  geom_line(aes(x=x, y=y), color="steelblue") +  
  labs(x="Smoothing Penalty", y="Generalized Cross Validation Score",  
       title="Generalized cross validated values of lambda") +  
  scale_x_log10() +  
  geom_vline(xintercept = lambdas[which(gcvs == min(gcvs))], color="coral")
```

The following code is used in Problem 1 part (b)

```

aicc = function(X=X, Y=Y, k=k, l){

  n = dim(X)[1]
  D = diag(c(rep(0, dim(X)[2]-k), rep(1, k)))
  H = X %%% solve(t(X) %%% X + l*D) %%% t(X)

  fhys = H %%% Y
  tr = sum(diag(H))

  r1 = log(mean((Y-fhys)^2))
  r2n = 2*(tr+1)
  r2d = n - tr - 2
  r = r1 + r2n/r2d
  return(r)
}

aicc_v = Vectorize(aicc, vectorize.args = c("l"))
lambdas = 10^(seq(-8, -3, length.out=100))
aiccs = aicc_v(X=X, Y=Y, l=lambdas, k=k)

aicc_df = data.frame(x=lambdas, y=aiccs)
plt_aicc = ggplot(data=aicc_df) +
  geom_line(aes(x=x, y=y), color="steelblue") +
  labs(x="Smoothing Penalty", y="Corrected AIC Score",
       title="Corrected AIC scores for smoothing parameter") +
  scale_x_log10() +
  geom_vline(xintercept = lambdas[which(aiccs == min(aiccs))], color="coral")

```

The following code is used in Problem 1 part (c)

```

sighat = function(Y=Y, X=X, k=k, l){
  D = diag(c(rep(0, dim(X)[2]-k), rep(1, k)))
  H = X %%% solve(t(X) %%% X + l*D) %%% t(X)
  fhys = H %%% Y
  r = mean((Y-fhys)^2)
  return(r)
}

er = function(X=X, Y=Y, k=k, l, lp){

  n = dim(X)[1]
  D = diag(c(rep(0, dim(X)[2]-k), rep(1, k)))
  H = X %%% solve(t(X) %%% X + l*D) %%% t(X)
  Hcv = X %%% solve(t(X) %%% X + lp*D) %%% t(X)

  fcvhys = Hcv %%% Y
  fhys = H %%% Y
  sig = sighat(X=X, Y=Y, k=k, l=lp)

  I = diag(dim(H)[1])
  r = (1/n)*(sum(((H-I)%%%fcvhys)^2) + sig*sum(diag(H%%t(H))))
  return(r)
}

erv = Vectorize(er, vectorize.args = c("l"))
lambdas = 10^(seq(-10, -3, length.out=100))
ers = erv(X=X, Y=Y, l=lambdas, k=k, lp=lcv)

er_df = data.frame(x=lambdas, y=ers)
plt_er = ggplot(data=er_df) +
  geom_line(aes(x=x, y=y), color="steelblue") +
  labs(x="Smoothing Penalty", y="Estimated Risk",
       title="Expected risk for smoothing parameter") +
  scale_x_log10() +
  geom_vline(xintercept = lambdas[which(ers == min(ers))], color="coral")

```

The following code is used in Problem 1 part (d)

```

## Parameterize
J = 6 # number of simulation parameterizations
M = 20 # number of simulations
L = 150 # number of lambdas to grid search
p = 3 # dimension of spline
k = 30 # number of knots
n = 200 # number of observations
a = 0
b = 1
knots = seq(a, b, length.out=k)
xs = ((1:n)-0.5)/n
fxs = f(xs)
lambdas = 10^seq(-10,0, length.out=L)

# Data storage objects

noise_raw_df = data.frame(x=xs, yt=fxs) # observations for plotting
noise_lam_df = data.frame( # best lambda for each simulation 1:M
  cv=rep(NA, J*M), gcv=rep(NA, J*M), aic=rep(NA, J*M), er=rep(NA, J*M),
  J=rep(NA, J*M))
noise_sco_df = data.frame( # score of the corresponding best lambda
  cv=rep(NA, J*M), gcv=rep(NA, J*M), aic=rep(NA, J*M), er=rep(NA, J*M),
  J=rep(NA, J*M))

D = diag(c(rep(0, p+1), rep(1, k)))
X1 = outer(xs, 0:p, "^")
X2 = outer(xs, knots, ">")*outer(xs, knots, "-")^p
X = cbind(X1, X2)

for(j in 1:J){

  for(i in 1:M){

    Y = fxs + rnorm(n, mean=0, sd=noise_sd(j))

    # calculate lambda using each method
    cvs = cvv(k=k, l=lambdas, X=X, Y=Y)
    gcvs = gcvv(k=k, l=lambdas, X=X, Y=Y)
    aiccs = aiccv(k=k, l=lambdas, X=X, Y=Y)

    bcv = min(cvs)
    cv_lam = lambdas[which(cvs == bcv)]

    ers = erv(k=k, l=lambdas, X=X, Y=Y, lp=cv_lam)

    bgcv = min(gcvs)
    gcv_lam = lambdas[which(gcvs == bgcv)]
    baicc = min(aiccs)
    aicc_lam = lambdas[which(aiccs == baicc)]
    ber = min(ers)
    er_lam = lambdas[which(ers == ber)]
  }
}

```

```

ix = (j-1)*M + i
noise_lam_df$J[ix] = j
noise_lam_df$cv[ix] = cv_lam
noise_lam_df$gcv[ix] = gcv_lam
noise_lam_df$aic[ix] = aicc_lam
noise_lam_df$er[ix] = er_lam

noise_sco_df$cv[ix] = bcv
noise_sco_df$gcv[ix] = bgcv
noise_sco_df$aic[ix] = baicc
noise_sco_df$er[ix] = ber

}

colname = paste("J=", j, sep="")
noise_raw_df[colname] = Y
}

## Plot results

plt_noise_obv = ggplot(data=melt(noise_raw_df, id=c("x", "yt"))) +
  geom_line(data=noise_raw_df, aes(x=x, y=yt), color="steelblue") +
  geom_point(aes(x=x, y=value, color=variable),
             alpha=0.7, size=1.5, shape=18) +
  facet_wrap(~variable) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Noise regime (J)"
  ) +
  labs(x="X", y="Y", title="Observations with different noise")

noise_lam_df_m = melt(noise_lam_df, id=c("J"))
noise_lam_df_m$J = as.factor(noise_lam_df_m$J)

ylims = boxplot.stats(noise_lam_df_m$value)$stats[c(1, 5)]

plt_noise_lams = ggplot(data=noise_lam_df_m) +
  geom_boxplot(aes(x=J, y=value, color=J), outlier.shape=NA) +
  coord_cartesian(ylim = ylims*1.05) +
  facet_wrap(~variable) +
  scale_y_log10() +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Noise regime (J)"
  ) +
  labs(x="Noise regime", y=TeX("Optimal  $\\lambda$ "),
       title=TeX("Stability of different methods for estimating  $\\lambda$ "))

```



```

## Parameterize
lambdas = 10^seq(-9,-4, length.out=L)

# Data storage objects
xts = ((1:n)-0.05)/n
yts = f(xts)
dens_raw_df = data.frame(x=xts, y=yts, j=rep(0, n)) # observations for plotting
dens_lam_df = data.frame( # best Lambda for each simulation 1:M
  cv=rep(NA, J*M), gcv=rep(NA, J*M), aic=rep(NA, J*M), er=rep(NA, J*M),
  J=rep(NA, J*M))
dens_sco_df = data.frame( # score of the corresponding best Lambda
  cv=rep(NA, J*M), gcv=rep(NA, J*M), aic=rep(NA, J*M), er=rep(NA, J*M),
  J=rep(NA, J*M))

D = diag(c(rep(0, p+1), rep(1, k)))

for(j in 1:J){

  for(i in 1:M){

    xs = qbeta(runif(n), (j+4)/5, (11-j)/5)
    fxs = f(xs)
    X1 = outer(xs, 0:p, "^")
    X2 = outer(xs, knots, ">")*outer(xs, knots, "-")^p
    X = cbind(X1, X2)

    Y = fxs + rnorm(n, mean=0, sd=0.1)

    # calculate lambda using each method
    cvs = cvv(k=k, l=lambdas, X=X, Y=Y)
    gcvs = gcvv(k=k, l=lambdas, X=X, Y=Y)
    aiccs = aiccv(k=k, l=lambdas, X=X, Y=Y)

    bcv = min(cvs)
    cv_lam = lambdas[which(cvs == bcv)]

    ers = erv(k=k, l=lambdas, X=X, Y=Y, lp=cv_lam)

    bgcv = min(gcvs)
    gcv_lam = lambdas[which(gcvs == bgcv)]
    baicc = min(aiccs)
    aicc_lam = lambdas[which(aiccs == baicc)]
    ber = min(ers)
    er_lam = lambdas[which(ers == ber)]

    ix = (j-1)*M + i
    dens_lam_df$J[ix] = j
    dens_lam_df$cv[ix] = cv_lam
    dens_lam_df$gcv[ix] = gcv_lam
    dens_lam_df$aic[ix] = aicc_lam
    dens_lam_df$er[ix] = er_lam
  }
}

```

```

dens_sco_df$cv[ix] = bcv
dens_sco_df$gcv[ix] = bgcv
dens_sco_df$aic[ix] = baicc
dens_sco_df$er[ix] = ber

}

tdf = data.frame(x=xs, y=Y, j=rep(j, n))
dens_raw_df = rbind(dens_raw_df, tdf)
}

dens_raw_df$j = as.factor(dens_raw_df$j)
drdf0 = subset(dens_raw_df, j == 0, select=c("x","y"))
plt_dens_obv = ggplot(data=drdf0) +
  geom_line(aes(x=x,y=y), color="steelblue") +
  geom_point(data=subset(dens_raw_df, j != 0), aes(x=x, y=y, color=j),
    shape=18, size=1.5) +
  facet_wrap(~j) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Support variance regime (J)"
  ) +
  labs(x="X", y="Y", title="Observations with variable support density")

dens_lam_df_m = melt(dens_lam_df, id=c("J"))
dens_lam_df_m$J = as.factor(dens_lam_df_m$J)
ylims = boxplot.stats(dens_lam_df_m$value)$stats[c(1, 5)]
plt_dens_lams = ggplot(data=dens_lam_df_m) +
  geom_boxplot(aes(x=J, y=value, color=J), outlier.shape=NA) +
  scale_y_log10() +
  coord_cartesian(ylim = ylims*1.05) +
  facet_wrap(~variable) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Support variance regime (J)"
  ) +
  labs(x="Support variance regime", y=TeX("Optimal  $\lambda$ "),
    title=TeX("Stability of different methods for estimating  $\lambda$ "))

```

```

## Parameterize
xs = ((1:n)-0.5)/n
lambdas = 10^seq(-11,-1, length.out=L)

ff = function(x, j){
  r = sqrt(x*(1-x))*sin((2*pi*(1+2^((9-4*j)/5)))/(x+2^((9-4*j)/5)))
  return(r)
}

# Data storage objects

sp_raw_df = data.frame(x=NA, yt=NA, ys=NA, j=NA)
sp_lam_df = data.frame(
  cv=rep(NA, J*M), gcv=rep(NA, J*M), aic=rep(NA, J*M), er=rep(NA, J*M),
  J=rep(NA, J*M))

D = diag(c(rep(0, p+1), rep(1, k)))
X1 = outer(xs, 0:p, "^")
X2 = outer(xs, knots, ">")*outer(xs, knots, "-")^p
X = cbind(X1, X2)

for(j in 1:J){

  fxs = ff(xs, j)

  for(i in 1:M){

    Y = fxs + rnorm(n, mean=0, sd=0.2)

    # calculate lambda using each method
    cvs = cvv(k=k, l=lambdas, X=X, Y=Y)
    gcvs = gcvv(k=k, l=lambdas, X=X, Y=Y)
    aiccs = aicc(k=k, l=lambdas, X=X, Y=Y)

    bcv = min(cvs)
    cv_lam = lambdas[which(cvs == bcv)]

    ers = erv(k=k, l=lambdas, X=X, Y=Y, lp=cv_lam)

    bgcv = min(gcvs)
    gcv_lam = lambdas[which(gcvs == bgcv)]
    baicc = min(aiccs)
    aicc_lam = lambdas[which(aiccs == baicc)]
    ber = min(ers)
    er_lam = lambdas[which(ers == ber)]

    ix = (j-1)*M + i
    sp_lam_df$J[ix] = j
    sp_lam_df$cv[ix] = cv_lam
    sp_lam_df$gcv[ix] = gcv_lam
    sp_lam_df$aic[ix] = aicc_lam
  }
}

```

```
    sp_lam_df$er[ix] = er_lam

  }

  tdf = data.frame(x=xs, yt=fxs, ys=Y, j=rep(j, n))
  sp_raw_df = rbind(sp_raw_df, tdf)
}

sp_raw_df = subset(sp_raw_df, ! is.na(j))
sp_raw_df$j = as.factor(sp_raw_df$j)
plt_sp_obv = ggplot(data=sp_raw_df) +
  geom_line(aes(x=x,y=yt), color="steelblue") +
  geom_point(aes(x=x, y=ys, color=j), shape=18, size=1.5) +
  facet_wrap(~j) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Function frequency regime (J)"
  ) +
  labs(x="X", y="Y",
    title="Observations generated by different target function frequencies")

sp_lam_df_m = melt(sp_lam_df, id=c("J"))
sp_lam_df_m$j = as.factor(sp_lam_df_m$j)
plt_sp_lams = ggplot(data=sp_lam_df_m) +
  geom_boxplot(aes(x=j, y=value, color=j)) +
  scale_y_log10() +
  facet_wrap(~variable) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Function frequency regime (J)"
  ) +
  labs(x="Target function frequencies", y=TeX("Optimal  $\lambda$ "),
    title=TeX("Stability of different methods for estimating  $\lambda$ "))
```

```

## Parameterize

xs = ((1:n)-0.5)/n
lambdas = 10^seq(-11,-1, length.out=L)

fv = function(x, j){
  r = (0.15*(1+0.4*(2*j-7)*(x-0.5)))^2
  return(sqrt(r))
}

# Data storage objects

v_raw_df = data.frame(x=NA, yt=NA, ys=NA, j=NA)
v_lam_df = data.frame(
  cv=rep(NA, J*M), gcv=rep(NA, J*M), aic=rep(NA, J*M), er=rep(NA, J*M),
  J=rep(NA, J*M))

D = diag(c(rep(0, p+1), rep(1, k)))
X1 = outer(xs, 0:p, "^")
X2 = outer(xs, knots, ">")*outer(xs, knots, "-")^p
X = cbind(X1, X2)
fxs = f(xs)

for(j in 1:J){
  for(i in 1:M){

    Y = fxs + fv(xs, j)*rnorm(n, mean=0, sd=1)

    # calculate lambda using each method
    cvs = cvv(k=k, l=lambdas, X=X, Y=Y)
    gcvs = gcvv(k=k, l=lambdas, X=X, Y=Y)
    aiccs = aiccv(k=k, l=lambdas, X=X, Y=Y)

    bcv = min(cvs)
    cv_lam = lambdas[which(cvs == bcv)]

    ers = erv(k=k, l=lambdas, X=X, Y=Y, lp=cv_lam)

    bgcv = min(gcvs)
    gcv_lam = lambdas[which(gcvs == bgcv)]
    baicc = min(aiccs)
    aicc_lam = lambdas[which(aiccs == baicc)]
    ber = min(ers)
    er_lam = lambdas[which(ers == ber)]

    ix = (j-1)*M + i
    sp_lam_df$J[ix] = j
    sp_lam_df$cv[ix] = cv_lam
    sp_lam_df$gcv[ix] = gcv_lam
  }
}

```

```

    sp_lam_df$aic[ix] = aicc_lam
    sp_lam_df$er[ix] = er_lam

  }

  tdf = data.frame(x=xs, yt=fxs, ys=Y, j=rep(j, n))
  v_raw_df = rbind(v_raw_df, tdf)
}

v_raw_df = subset(v_raw_df, ! is.na(j))
v_raw_df$j = as.factor(v_raw_df$j)
plt_v_obv = ggplot(data=v_raw_df) +
  geom_line(aes(x=x,y=yt), color="steelblue") +
  geom_point(aes(x=x, y=ys, color=j), shape=18, size=1.5) +
  facet_wrap(~j) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Variance regime (J)"
  ) +
  labs(x="X", y="Y",
    title="Observations generated by different variance structures")

v_lam_df_m = melt(sp_lam_df, id=c("J"))
v_lam_df_m$j = as.factor(v_lam_df_m$j)
plt_v_lams = ggplot(data=v_lam_df_m) +
  geom_boxplot(aes(x=j, y=value, color=j)) +
  scale_y_log10() +
  facet_wrap(~variable) +
  scale_colour_discrete(
    labels=as.character(1:J),
    name="Variance regime (J)"
  ) +
  labs(x="Variance structures", y=TeX("Optimal  $\lambda$ "),
    title=TeX("Stability of different methods for estimating  $\lambda$ "))

```